

Im following this tutorial: <https://learnopencv.com/fine-tuning-yolov7-on-custom-dataset/>

```
1 import os
2 import torch
3
4 if torch.cuda.is_available():
5     print("CUDA is available! 🍌")
6     print(f"CUDA version: {torch.version.cuda}")
7     print(f"Number of GPUs: {torch.cuda.device_count()}")
8 else:
9     print("CUDA is not available. 😞")
10
11 !nvidia-smi
```

```
🔗 CUDA is available! 🍌
CUDA version: 12.1
Number of GPUs: 1
Sat Jan 25 09:45:48 2025
```

NVIDIA-SMI 535.104.05			Driver Version: 535.104.05			CUDA Version: 12.2		
GPU Fan	Name Temp	Perf	Persistence-M Pwr:Usage/Cap	Bus-Id	Disp.A Memory-Usage	Volatile GPU-Util	Uncorr. Compute M.	ECC MIG M.
0 N/A	Tesla 34C	T4 P8	Off 10W / 70W	00000000:00:04.0	Off 3MiB / 15360MiB	0%	Default	0 N/A
+-----+								
+-----+								
Processes:								
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage		
	ID	ID						
=====								
No running processes found								

```
1 # Download the dataset. ifnot already there.
2 if not os.path.exists('pothole_dataset.zip'):
3     !wget https://learnopencv.s3.us-west-2.amazonaws.com/pothole_dataset.zip
4
```

Download dataset, extract code and view direcotry tree

```
1 # Extract the dataset.
2 if not os.path.exists('pothole_dataset'):
3     !unzip -q pothole_dataset.zip
```

```
1 !ls -R pothole_dataset
```

🔗 [Show hidden output](#)

We need the above output of the direcotry tree to look something like this.

NOTE TO SELF: Check later if this is an actual requirement of training a YOLOv7 model.

```
pothole_dataset/
├── images
│   ├── test [118 entries exceeds filelimit, not opening dir]
│   ├── train [1265 entries exceeds filelimit, not opening dir]
│   └── valid [401 entries exceeds filelimit, not opening dir]
└── labels
    ├── test [118 entries exceeds filelimit, not opening dir]
    ├── train [1265 entries exceeds filelimit, not opening dir]
    └── valid [401 entries exceeds filelimit, not opening dir]
```

Below are some interesting yaps about how the boundary boxes are labeled along with the images.

Each image comes along with a text file that contains the annotations of the boundary boxes for that image.

```
0 0.5497282608695652 0.5119565217391304 0.017934782608695653 0.005072463768115942
0 0.41032608695652173 0.5253623188405797 0.025 0.005797101449275362
```

```
0 0.30842391304347827 0.5282608695652173 0.014673913043478261 0.005797101449275362
0 0.1654891304347826 0.5224637681159421 0.027717391304347826 0.005797101449275362
0 0.10163043478260869 0.5286231884057971 0.01956521739130435 0.006521739130434782
0 0.07907608695652174 0.5293478260869565 0.01576086956521739 0.007971014492753623
```

- First number: Class (category)
- Next four floats: <x_center, y_center, width, height>

All coordinates are noramlised to the image.

✓ CLONE YOLO!!!

```
1 if not os.path.exists('yolov7'):
2     !git clone https://github.com/WongKinYiu/yolov7.git
3
4 %cd yolov7
5
6 !pip install -r requirements.txt
```

 [Show hidden output](#)

Clone if not already there then cd into it and install requirements.

✓ FEED THE CLONE YALM!!!


Fuck the fluff

To train YOLO models you need configured .yalm files.

lets see how thats done

"This .yaml file contains the paths to the image sets, the number of classes, and the name of the classes."

```
1 %%writefile data/pothole.yaml
2 train: ../pothole_dataset/images/train
3 val: ../pothole_dataset/images/valid
4 test: ../pothole_dataset/images/test
5
6 # Classes
7 nc: 1 # number of classes
8 names: ['pothole'] # class names
```

 Writing data/pothole.yaml

Remember this is still in the yolov7 directory.

Right now i cant see it tho...hmm weird, strange behaviour...

✓ Lets start small: Tiny YOLOv7 Fixed Resolution Training

OOhhhh 6 million parameters, ahhh woowwww so impressive..... idk, actually? is it?

native base resolution: 640x640

1. First get the model:

```
1 # Download the Tiny model weights.
2 !wget https://github.com/WongKinYiu/yolov7/releases/download/v0.1/yolov7-tiny.pt
```

 [Show hidden output](#)

This has been pre-trained on the COCO dataset. Ill link that later on... if i remember.

YOOO, little side note, i found the crab mode feature under the misc settings. and this combo feature is fiiiiirrrrrreeeeee my high score right now is like 250.

fun

2. Configuration: they are vague on what were configuring but config files for the model can be found here: yolov7/cfg/training/

The high score is now 600...

Anyway, we need to change the configuration file current present there with the below code. Actually, just for comparison, here is the current code:

```
# parameters
nc: 80 # number of classes
depth_multiple: 1.0 # model depth multiple
width_multiple: 1.0 # layer channel multiple

# anchors
anchors:
  - [10,13, 16,30, 33,23] # P3/8
  - [30,61, 62,45, 59,119] # P4/16
  - [116,90, 156,198, 373,326] # P5/32

# yolov7-tiny backbone
backbone:
  # [from, number, module, args] c2, k=1, s=1, p=None, g=1, act=True
  [[-1, 1, Conv, [32, 3, 2, None, 1, nn.LeakyReLU(0.1)]], # 0-P1/2

    [-1, 1, Conv, [64, 3, 2, None, 1, nn.LeakyReLU(0.1)]], # 1-P2/4

    [-1, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-2, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [[-1, -2, -3, -4], 1, Concat, [1]],
    [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 7

    [-1, 1, MP, []], # 8-P3/8
    [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-2, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [[-1, -2, -3, -4], 1, Concat, [1]],
    [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 14

    [-1, 1, MP, []], # 15-P4/16
    [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-2, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [[-1, -2, -3, -4], 1, Concat, [1]],
    [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 21

    [-1, 1, MP, []], # 22-P5/32
    [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-2, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [256, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, Conv, [256, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
    [[-1, -2, -3, -4], 1, Concat, [1]],
    [-1, 1, Conv, [512, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 28
  ]

# yolov7-tiny head
head:
  [[-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-2, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [-1, 1, SP, [5]],
    [-2, 1, SP, [9]],
    [-3, 1, SP, [13]],
    [[-1, -2, -3, -4], 1, Concat, [1]],
    [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
    [[-1, -7], 1, Concat, [1]],
    [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 37

    [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
```

```

[-1, 1, nn.Upsample, [None, 2, 'nearest']],
[21, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # route backbone P4
[[-1, -2], 1, Concat, [1]],

[-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-2, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[[-1, -2, -3, -4], 1, Concat, [1]],
[-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 47

[-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, nn.Upsample, [None, 2, 'nearest']],
[14, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # route backbone P3
[[-1, -2], 1, Concat, [1]],

[-1, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-2, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[[-1, -2, -3, -4], 1, Concat, [1]],
[-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 57

[-1, 1, Conv, [128, 3, 2, None, 1, nn.LeakyReLU(0.1)]],
[[-1, 47], 1, Concat, [1]],

[-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-2, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[[-1, -2, -3, -4], 1, Concat, [1]],
[-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 65

[-1, 1, Conv, [256, 3, 2, None, 1, nn.LeakyReLU(0.1)]],
[[-1, 37], 1, Concat, [1]],

[-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-2, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[[-1, -2, -3, -4], 1, Concat, [1]],
[-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 73

[57, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[65, 1, Conv, [256, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
[73, 1, Conv, [512, 3, 1, None, 1, nn.LeakyReLU(0.1)]],

[[74,75,76], 1, IDetect, [nc, anchors]], # Detect(P3, P4, P5)
]

1 %%writefile cfg/training/yolov7_pothole-tiny.yaml
2 # parameters
3 nc: 1 # number of classes
4 depth_multiple: 1.0 # model depth multiple
5 width_multiple: 1.0 # layer channel multiple
6
7 # anchors
8 anchors:
9   - [10,13, 16,30, 33,23] # P3/8
10  - [30,61, 62,45, 59,119] # P4/16
11  - [116,90, 156,198, 373,326] # P5/32
12
13 # yolov7-tiny backbone
14 backbone:
15   # [from, number, module, args] c2, k=1, s=1, p=None, g=1, act=True
16   [[-1, 1, Conv, [32, 3, 2, None, 1, nn.LeakyReLU(0.1)]], # 0-P1/2
17
18   [-1, 1, Conv, [64, 3, 2, None, 1, nn.LeakyReLU(0.1)]], # 1-P2/4
19
20   [-1, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
21   [-2, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
22   [-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],

```

```

23 [-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
24 [[-1, -2, -3, -4], 1, Concat, [1]],
25 [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 7
26
27 [-1, 1, MP, []], # 8-P3/8
28 [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
29 [-2, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
30 [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
31 [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
32 [[-1, -2, -3, -4], 1, Concat, [1]],
33 [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 14
34
35 [-1, 1, MP, []], # 15-P4/16
36 [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
37 [-2, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
38 [-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
39 [-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
40 [[-1, -2, -3, -4], 1, Concat, [1]],
41 [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 21
42
43 [-1, 1, MP, []], # 22-P5/32
44 [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
45 [-2, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
46 [-1, 1, Conv, [256, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
47 [-1, 1, Conv, [256, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
48 [[-1, -2, -3, -4], 1, Concat, [1]],
49 [-1, 1, Conv, [512, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 28
50 ]
51
52 # yolov7-tiny head
53 head:
54 [[-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
55 [-2, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
56 [-1, 1, SP, [5]],
57 [-2, 1, SP, [9]],
58 [-3, 1, SP, [13]],
59 [[-1, -2, -3, -4], 1, Concat, [1]],
60 [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
61 [[-1, -7], 1, Concat, [1]],
62 [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 37
63
64 [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
65 [-1, 1, nn.Upsample, [None, 2, 'nearest']],
66 [21, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # route backbone P4
67 [[-1, -2], 1, Concat, [1]],
68
69 [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
70 [-2, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
71 [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
72 [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
73 [[-1, -2, -3, -4], 1, Concat, [1]],
74 [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 47
75
76 [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
77 [-1, 1, nn.Upsample, [None, 2, 'nearest']],
78 [14, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # route backbone P3
79 [[-1, -2], 1, Concat, [1]],
80
81 [-1, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
82 [-2, 1, Conv, [32, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
83 [-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
84 [-1, 1, Conv, [32, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
85 [[-1, -2, -3, -4], 1, Concat, [1]],
86 [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 57
87
88 [-1, 1, Conv, [128, 3, 2, None, 1, nn.LeakyReLU(0.1)]],
89 [[-1, 47], 1, Concat, [1]],
90
91 [-1, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
92 [-2, 1, Conv, [64, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
93 [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
94 [-1, 1, Conv, [64, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
95 [[-1, -2, -3, -4], 1, Concat, [1]],
96 [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 65
97
98 [-1, 1, Conv, [256, 3, 2, None, 1, nn.LeakyReLU(0.1)]],
99 [[-1, 37], 1, Concat, [1]],
100
101 [-1, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
102 [-2, 1, Conv, [128, 1, 1, None, 1, nn.LeakyReLU(0.1)]],
103 [-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
104 [-1, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],

```

```

105 [[-1, -2, -3, -4], 1, Concat, [1]],
106 [-1, 1, Conv, [256, 1, 1, None, 1, nn.LeakyReLU(0.1)]], # 73
107
108 [57, 1, Conv, [128, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
109 [65, 1, Conv, [256, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
110 [73, 1, Conv, [512, 3, 1, None, 1, nn.LeakyReLU(0.1)]],
111
112 [[74,75,76], 1, IDetect, [nc, anchors]], # Detect(P3, P4, P5)
113 ]

```

✎ Writing cfg/training/yolov7_pothole-tiny.yaml

Alright, here we have some model architecture, this is the good stuff. NICCEE

We needed to change the number of classes to 1.

✓ Let the Training BEGIN!

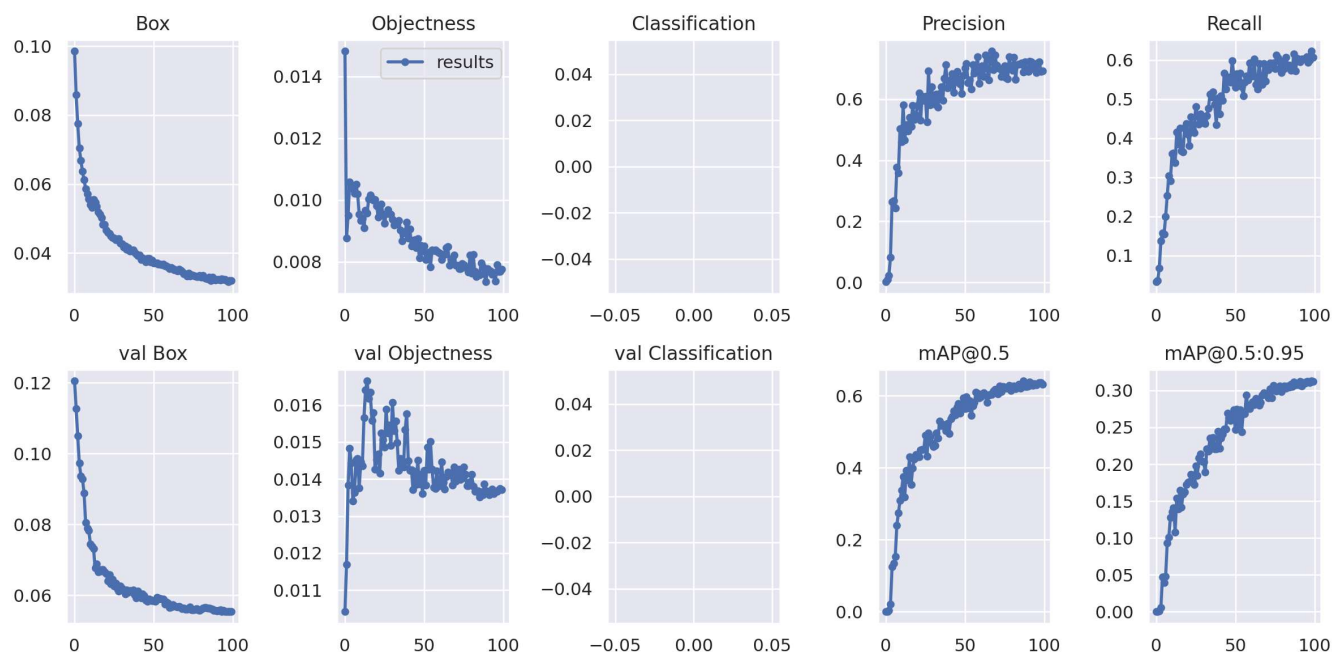
```

1 !python train.py --epochs 100 --workers 8 --device 0 --batch-size 32 \
2 --data data/pothole.yaml --img 640 640 --cfg cfg/training/yolov7_pothole-tiny.yaml \
3 --weights 'yolov7-tiny.pt' --name yolov7_tiny_pothole_fixed_res --hyp data/hyp.scratch.tiny.yaml

```

✎ Show hidden output

For a small model this thing does relatively fine. Here are some results:



Now we get to run a test:

```
1 !python test.py --weights runs/train/yolov7_tiny_pothole_fixed_res/weights/best.pt --task test --data data/pothole.yaml
```

✎ Show hidden output

✓ Tiny Multi-Resolution Training

Right, our next step is doing multi resolution training.

Here is what is going to be happening:

1. we set a base resolution like 640x640
2. Training then happens in batches
3. Images in the batches can be +50% of base resolution
4. so between 320x320 and 1280x1280 (according to learnopencv) 640+50%=960... not 1280.
5. Results in a more robust model that deals with small details
6. Need to train for longer as data is more diverse

```

1 !python train.py --epochs 100 --workers 4 --device 0 --batch-size 32 \
2 --data data/pothole.yaml --img 640 640 --cfg cfg/training/yolov7_pothole-tiny.yaml \
3 --weights 'yolov7-tiny.pt' --name yolov7_tiny_pothole_multi_res --hyp data/hyp.scratch.tiny.yaml \
4 --multi-scale

```

```

2/99    0.776G    0.079 0.009544    0 0.08855    58    352: 100% 40/40 [01:11<00:00, 1.79s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:07<00:00, 1.04s/it]
      all      401      1034    0.00989    0.0571    0.00131    0.000252

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
3/99    1.16G    0.07209 0.01257    0 0.08466    44    928: 100% 40/40 [01:14<00:00, 1.87s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:04<00:00, 1.40it/s]
      all      401      1034    0.0655    0.0687    0.00974    0.00209

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
4/99    0.795G    0.06639 0.01319    0 0.07958    88    736: 100% 40/40 [01:09<00:00, 1.74s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:07<00:00, 1.04s/it]
      all      401      1034    0.174    0.113    0.0497    0.0162

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
5/99    11.7G    0.06205 0.01245    0 0.07451    81    352: 100% 40/40 [01:08<00:00, 1.71s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:08<00:00, 1.19s/it]
      all      401      1034    0.244    0.158    0.102    0.0364

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
6/99    11.7G    0.05889 0.01134    0 0.07023    46    352: 100% 40/40 [01:08<00:00, 1.71s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:06<00:00, 1.05it/s]
      all      401      1034    0.255    0.176    0.14    0.0511

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
7/99    1.26G    0.05662 0.01225    0 0.06887    68    896: 100% 40/40 [01:12<00:00, 1.81s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:06<00:00, 1.07it/s]
      all      401      1034    0.372    0.215    0.189    0.0688

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
8/99    0.468G    0.05545 0.01275    0 0.06819    54    512: 100% 40/40 [01:11<00:00, 1.80s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:05<00:00, 1.26it/s]
      all      401      1034    0.402    0.283    0.277    0.102

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
9/99    1.47G    0.05287 0.01223    0 0.0651    68    960: 100% 40/40 [01:14<00:00, 1.86s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:06<00:00, 1.08it/s]
      all      401      1034    0.477    0.262    0.26    0.0909

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
10/99   0.575G    0.05168 0.01274    0 0.06442    80    384: 100% 40/40 [01:09<00:00, 1.75s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:08<00:00, 1.20s/it]
      all      401      1034    0.446    0.35    0.325    0.13

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
11/99   0.82G    0.05071 0.01152    0 0.06223    37    640: 100% 40/40 [01:09<00:00, 1.75s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:06<00:00, 1.12it/s]
      all      401      1034    0.474    0.325    0.323    0.132

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
12/99   11.9G    0.05186 0.01097    0 0.06283    76    736: 100% 40/40 [01:05<00:00, 1.65s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 100% 7/7 [00:09<00:00, 1.35s/it]
      all      401      1034    0.363    0.32    0.255    0.0884

Epoch   gpu_mem    box    obj    cls    total    labels    img_size
13/99   0.778G    0.0503 0.0102    0 0.0605    51    672: 100% 40/40 [01:13<00:00, 1.84s/it]
      Class    Images    Labels    P    R    mAP@.5    mAP@.5:.95: 0% 0/7 [00:00<?, ?it/s]

```

```

1 !python train.py \
2 --weights last.pt \
3 --cfg cfg/training/yolov7-tiny.yaml \
4 --data data.yaml \
5 --epochs 100 --resume
6

```

python3: can't open file '/content/train.py': [Errno 2] No such file or directory

